



Kubernetes for Application Developers

Why We Use It

Intermediate Kubernetes — Module 1 of 13



Module Objectives

- Articulate the developer value proposition of Kubernetes
- Trace the journey from source code to running cluster workload
- Define the developer-platform contract and responsibilities
- Review core abstractions (Pod, Deployment, Service) through a developer lens
- Compare declarative vs. imperative approaches
- Connect 12-factor app principles to Kubernetes patterns
- Preview the full 13-module course roadmap



The Problem We're Solving

Why developers needed something better

Before Kubernetes — Developer Pain Points

Deployment Challenges

- "Works on my machine" syndrome
- Manual server provisioning tickets
- Environment drift between dev/staging/prod
- Multi-hour deployment windows

Operational Burden

- Developers on-call for infrastructure issues
- Scaling requires capacity planning meetings
- Rollback = panic + SSH + hope
- No self-service for infrastructure needs

✓ After Kubernetes — Developer Gains

Deployment Gains

- Containers guarantee environment parity
- Self-service deployments via `kubectl apply`
- Same manifest works dev → staging → prod
- Rolling updates with zero downtime

Operational Gains

- Platform team owns infrastructure
- Auto-scaling based on actual metrics
- Rollback = `kubectl rollout undo`
- Declarative config = version-controlled infra

Before vs. After — Side by Side

Concern	Before K8s	With K8s
Deploy a new version	SSH, scripts, prayer	<code>kubectl apply -f deployment.yaml</code>
Scale to handle load	File a ticket, wait days	HPA scales automatically
Rollback a bad release	Restore from backup, redeploy	<code>kubectl rollout undo</code>
Service discovery	Hardcoded IPs, load balancer configs	DNS-based: <code><service>.<namespace>.svc.cluster.local</code>
Environment parity	Snowflake servers	Same container image everywhere
Health monitoring	External monitoring + manual restart	Liveness/readiness probes + auto-restart



The Developer Workflow

From code to cluster



Code to Cluster — The Journey

The Developer Pipeline

1. **Write Code** — Application logic, tests, Dockerfile
2. **Build Image** — `docker build -t myapp:v1.2.3 .`
3. **Push to Registry** — `docker push ecr-repo/myapp:v1.2.3`
4. **Declare Desired State** — Write/update Kubernetes manifests
5. **Apply to Cluster** — `kubectl apply -f k8s/`
6. **Kubernetes Reconciles** — Controller loop makes it so
7. **Observe & Iterate** — Logs, metrics, traces



What Developers Own

Application Artifacts

- Application source code and tests
- Dockerfile / container definition
- Application configuration

Kubernetes Manifests

- Deployment and Service definitions
- ConfigMaps and Secrets
- Resource requests, limits, and health probes

Key Insight: Developers define *what* they need. The platform provides *how* it's delivered.

The Developer-Platform Contract

Clear boundaries, clear responsibilities

Developer vs. Platform Responsibilities

Developer Responsibility	Platform Responsibility
Container image (Dockerfile)	Container runtime (containerd)
Resource requests & limits	Node capacity & scheduling
Health check endpoints	Health check execution & remediation
Deployment strategy choice	Rolling update orchestration
Service port definitions	Network routing & load balancing
PVC requests (size, access mode)	Storage provisioning (EBS, EFS)
Scaling thresholds (HPA config)	Cluster autoscaler, node pools

The Kubernetes API — Your Contract Interface

The API Server is the single point of interaction

kubectl, CI/CD pipelines, controllers, and the scheduler all speak the same API.

Why this matters: The API is versioned and consistent. Your manifests work on future clusters thanks to API compatibility guarantees.

Self-Service Infrastructure

Infrastructure requests become API calls -- no tickets required:

I need...

- A load balancer
- Persistent storage
- A cron job
- DNS entry and TLS termination
- Horizontal scaling

I write...

- Service type: LoadBalancer
- PersistentVolumeClaim
- CronJob
- Service + Ingress with TLS
- HorizontalPodAutoscaler



Core Abstractions Review

Pods, Deployments, and Services — the developer lens



Abstraction-to-Need Mapping

Developer Need	K8s Abstraction	What It Provides
Run my container	Pod	Smallest deployable unit, shared network/storage
Keep N replicas running	Deployment / ReplicaSet	Desired state, rolling updates, rollback
Stable network endpoint	Service	Load balancing, DNS, stable IP
External traffic routing	Ingress / Gateway	HTTP routing, TLS, path-based routing
Configuration management	ConfigMap / Secret	Externalized config, decoupled from image
Persistent data	PVC / StorageClass	Dynamic provisioning, lifecycle management
Run-to-completion tasks	Job / CronJob	Batch processing, scheduled tasks



Pods — The Atomic Unit

- Smallest deployable unit in K8s
- One or more containers sharing network & storage
- Every container in a pod shares `localhost`
- Pods are **ephemeral** — they can be killed and rescheduled
- You rarely create pods directly — use Deployments

Deployments — Desired State Management

What Deployments provide:

- **Desired state** — declare replica count and Kubernetes reconciles continuously
- **Rolling updates** — new pods start before old ones terminate
(`maxUnavailable: 0`)
- **Rollback** — instantly revert to a previous revision with `kubectl rollout undo`
- **Self-healing** — failed pods are automatically replaced to match the desired count

Services — Stable Networking

Why Services exist:

- Pods have dynamic IPs — they change on restart
- Services provide a **stable** virtual IP
- Built-in load balancing across pod replicas
- DNS-based discovery — no hardcoded IPs
- Selector connects Service to Pods via labels

Think of it as: A stable address that always routes to healthy pods.



Declarative vs. Imperative

The most important mental shift

Two Ways to Interact with Kubernetes

Imperative (Not Recommended)

- Step-by-step commands:
create, scale, set image
- No record of intended state
- Hard to audit or reproduce

Tells Kubernetes *what to do* step by step.

Declarative (Best Practice)

- Declare desired state in a manifest
- Update by editing the YAML and re-applying
- Everything lives in version control

Tells Kubernetes *what you want* — it figures out how.

Why Declarative Wins



Auditability

- Git history and PR reviews for every change
- Blame shows who changed what



Reproducibility

- Same YAML = same result everywhere
- Disaster recovery is a `kubectl apply`



Automation

- GitOps workflows (ArgoCD, Flux)
- CI/CD pipelines with drift detection



Environment Parity

Same image, same behavior, everywhere



Environment Parity with Kubernetes

The Promise

The same container image runs identically in every environment.
Configuration is injected, not baked in.

What stays the same

- Container image (immutable artifact)
- Application code, dependencies, and runtime
- OS-level packages

What changes per environment

- ConfigMaps and Secrets
- Resource requests/limits and replica counts
- Ingress rules / hostnames



12-Factor App Principles

How Kubernetes embodies modern application design



12-Factor Principles & Kubernetes

Factor	Principle	Kubernetes Implementation
I	Codebase	One repo, one container image
II	Dependencies	Bundled in container image
III	Config	ConfigMaps, Secrets, env vars
IV	Backing Services	Services, ExternalName, endpoints
V	Build, Release, Run	Image build → tag → Deployment
VI	Processes	Stateless pods, external state



12-Factor Principles & Kubernetes (cont.)

Factor	Principle	Kubernetes Implementation
VII	Port Binding	containerPort + Service
VIII	Concurrency	Horizontal scaling (replicas, HPA)
IX	Disposability	Fast startup, graceful shutdown (preStop)
X	Dev/Prod Parity	Same image + Kustomize overlays
XI	Logs	Write to stdout/stderr, collected by platform
XII	Admin Processes	Jobs, CronJobs, <code>kubectl exec</code>

Design implication: Applications that follow these principles are "Kubernetes-native" — they leverage the platform fully and avoid fighting it.

Factor IX Deep Dive — Disposability

Graceful shutdown sequence:

- Pod is removed from Service endpoints — no new traffic is routed
- **preStop hook** runs first — a brief drain allows in-flight requests to complete
- **SIGTERM** is sent to the container process — the app should finish work and exit
- After `terminationGracePeriodSeconds`, Kubernetes sends **SIGKILL** to force termination



Kubernetes Architecture

What's under the hood — and why developers should care



Control Plane vs. Worker Nodes



Control Plane (AWS-managed on EKS)

- **API Server** — the front door; everything talks to it
- **etcd** — key-value store, the cluster's source of truth
- **Scheduler** — places Pods onto nodes
- **Controller Manager** — reconcile loops driving actual → desired



Worker Nodes (your EC2 instances)

- **kubelet** — runs and watches your Pods
- **kube-proxy** — programs Service networking on the node
- **container runtime** — containerd; pulls images, runs containers

Developer takeaway: you only ever talk to the *API server* (via `kubectl / manifests`); controllers and the kubelet make your declared state real

The Pluggable Interfaces

The core doesn't hard-code runtime, networking, or storage — each is a swappable interface:

Interface	Plugs in	On EKS
CRI — Container Runtime	How containers run	containerd
CNI — Container Network	Pod networking & IPs	AWS VPC CNI + Calico
CSI — Container Storage	Volume provisioning	EBS CSI driver

Why it matters: the gp3 StorageClass, Pod IPs, and NetworkPolicies you'll use in later labs are delivered by these plug-in drivers — not the Kubernetes core.

Kubernetes on EKS

Your enterprise platform

Amazon EKS — What It Means for Developers

AWS Manages

- Control plane (API server, etcd, scheduler, controllers)
- HA and multi-AZ availability
- Version upgrades and certificates

Your Team Manages

- Worker node groups (EC2 or Fargate)
- Application deployments and configuration
- Namespaces, RBAC, and add-ons

EKS advantage: The control plane is someone else's problem. You focus on workloads.



Enterprise Architecture Patterns

Mapping N-tier applications to Kubernetes primitives

Common N-Tier Application on Kubernetes

Frontend

- **Deployment** — stateless replicas
- **Service** — internal routing
- **Ingress** — external traffic + TLS

API / Backend

- **Deployment** — horizontally scaled
- **Service** — DNS discovery
- **ConfigMap / Secret** — config

Data Layer

- **StatefulSet + PVC** — stable storage
- **Headless Service** — pod DNS
- **Cache** — Redis / Memcached

Every tier maps to a small set of primitives -- master these and you can model any architecture.



Enterprise Architecture on EKS

- **DNS-critical services** — multi-AZ scheduling + PodDisruptionBudgets
- **Dual Ingress** — separate internal vs. external traffic
- **GitOps via FluxCD** — Git as single source of truth
- **Vault + ESO** — secrets management with Kubernetes auth
- **ECR** — scanned container images
- **Observability** — Prometheus, Grafana, Splunk

Each pattern is covered hands-on in Modules 5-13.



Course Roadmap

#	Module	Focus
1	Why Kubernetes	Core abstractions, EKS
2	Scaling	HPA, VPA, Cluster Autoscaler
3	Storage	PVCs, StatefulSets
4	Services	Service types, DNS
5	Config & Secrets	ConfigMaps, Secrets, Vault
6	Ingress & Egress	Ingress, Gateway API, TLS
7	RBAC & Security	Roles, PSA, IRSA
8	Network Policies	Pod-to-pod traffic control



Course Roadmap — Modules 9-13

#	Module	Focus
9	Observability	Logging (Splunk), metrics, Grafana, alerting
10	Health Probes & Debugging	Liveness, readiness, failure states, troubleshooting
11	Deployment Strategies	Rolling, blue-green, canary, progressive delivery
12	Helm & Kustomize	Charts, templating, release management
13	CI/CD & GitOps	Jenkins, ArgoCD, FluxCD

Lab 1: Exploring Your Kubernetes Cluster

Hands-on: Connect, explore, deploy

- Connect to the EKS cluster using `aws eks update-kubeconfig`
- Explore cluster resources with `kubectl get` and `kubectl describe`
- Deploy a sample application with `kubectl create deployment`
- Expose it with a Service and verify connectivity
- Scale the Deployment up and down and observe replica changes
- Examine pod lifecycle events and logs

Duration: ~45 minutes



Module 1 — Key Takeaways

- **Kubernetes is a developer platform** — it transforms infrastructure requests into self-service API calls
- **The developer-platform contract** — you define what you need; the platform delivers how
- **Core abstractions map to developer needs** — Pod (run), Deployment (manage), Service (connect)
- **Declarative > imperative** — YAML in Git is your source of truth
- **Environment parity via containers** — same image, different config per environment
- **12-factor apps thrive on K8s** — the platform was designed for this pattern
- **EKS simplifies operations** — AWS manages the control plane, you manage workloads

N-tier apps map cleanly to K8s primitives — Deployment, Service

? Questions & Discussion

What aspects of your current workflow would you most like
Kubernetes to improve?

Up Next: Module 2 — Scaling Your Applications